


A DATA-SCIENCE TEAM, ON DEMAND

Hey OctOpus,
Connect my data warehouse.
Transform my sales data.
Show my top customers.
Build a *dashboard*.
Forecast next quarter's revenue.

Point OctOpus at your source – *Databricks, Snowflake, BigQuery, Postgres, a file* – and have a normal conversation. It connects, cleans, analyzes, builds live dashboards, and trains & deploys forecasting models – the work of a whole data-science team, without the hiring queue or the backlog. **Your analysts, consultants, and operators ship data science themselves** – in hours, not quarters.



CONNECTS DATABRICKS · SNOWFLAKE · BIGQUERY DEPLOYS IN YOUR ENVIRONMENT

TIME TO VALUE HOURS, NOT QUARTERS PROVEN LIVE WITH FORTUNE 500 TEAMS RTNERS & PROGRAMS



THE PROBLEM

Reliable models, *on demand*, for decisions that cannot wait.

Getting from a raw warehouse to a decision still means exports, SQL, notebooks, packages, validation, retries, and deployment – work that demands data-science talent that's scarce, expensive, and already stretched thin. The whole loop is too technical and too manual.



01

Obsolete AutoML

Benchmarks models. Still needs an expert to frame the task, iterate, debug, and deploy.

NO REASONING

NO CONTEXT

NO SELF-DEBUG

NO DEPLOYMENT



02

Pain / Complexity

Reliable models still mean notebooks, coding tools, packages, MLOps, and expert back-and-forth.

NOTEBOOKS

CODING TOOLS

PACKAGES

MLOPS

EXPERTISE

TIME



03

Reliability and speed

OctOpus is AI-native. It does the reasoning and turns one question into a reliable deployed model.

DATA SCIENTISTS

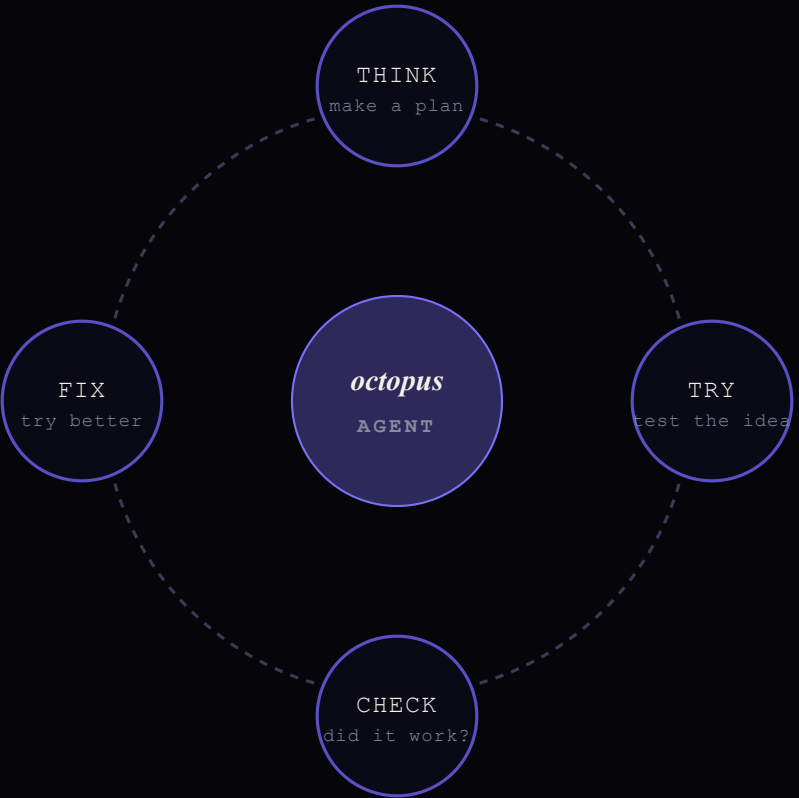
Build and deploy
reliable models faster
and more frequently.

BUSINESS USERS

“Forecast X per Y for
Z.” Ask on real data,
get visuals and
dashboards, then build
and deploy the model.

ONE CONVERSATION

From your *raw warehouse* to a clean answer, a live dashboard, and a deployed model.
You just talk — and get senior data-science work, on demand.



- 01 →

“Connect to my sales data in Databricks.”

OctOpus connects to your source – Databricks, Snowflake, BigQuery, Postgres, or a file – and loads it. No export, no query.
- 02 →

“Clean it up and show me my top customers.”

It transforms the data and answers – your top customers, ranked, with the numbers behind it.
- 03 →

“Build me a dashboard from this.”

It generates a live, interactive dashboard of your sales – no BI tool, no setup.
- 04 →

“Forecast next quarter's revenue — train a model.”

It builds, tests, and deploys a forecasting model, then has a prediction. One conversation, raw data to deployed model.



SEE IT WORK

Watch one conversation go from raw data to a deployed model.

OPEN THE DEMO

`how-to-use.html ↗ - how-to-use.html``www.octopus.dev ↗ - https://www.octopus.dev/`

ALREADY LIVE WITH REAL TEAMS

Real teams. Real data. *Real problems solved.*

Your team doesn't just want a prediction. They want OctOpus to do the whole job — analyze the data, build the model, fill in the reports, and hand it back ready to use. That's exactly what it does today.

WHAT TEAMS GET OUT OF IT

- Answers from raw warehouse data in hours
- Live dashboards and deployed models, no MLOps
- Reports filled in automatically, ready to send



200+

PROFESSIONALS USING IT

Analysts, researchers, consultants, big-company teams

9

ENTERPRISE POCS RUNNING

Logistics, lean management, deeptech, oil & gas, SaaS · US, France, Sweden, Africa

4 / 9

POCS CONVERTED TO PAID

The rest in late-stage rollout · PoCs become contracts

Hours

FROM DATA TO MODEL

The work that takes a hire or AutoML weeks to months

INBOUND CONVERSATIONS · Q2 2026

IBM

Volvo Group

SKF

Ericsson

AstraZeneca

Actemium

Active inbound discovery calls with Fortune 500 enterprises. No outbound sales — they came to us.

A REAL EXAMPLE

A factory client makes bottles. OctOpus analyzed their quality data, found the defects, and filled in *10 sections of their existing client report* — automatically. Their team just had to sign and send.

Writing code is easy. *Specialized knowledge and reasoning* are what you're really buying.

Claude, Codex, and Cursor can all write code. The hard part is knowing what to test, what to fix, what to trust, and what to ship for your specific forecasting problem. OctOpus does that reasoning for you – and gets sharper on your data with every run.

	AUTOML TOOLS DataRobot, H2O,	CODING AGENTS Claude, Codex, Cursor,	HIRE A DATA SCIENTIST Strong, but slow and	OCTOPUS AI-native data science for specialists and operators Reason, analyze, train, report One system across analysis and data science
What they do best	Benchmark models Good for narrow structured ML flows	Write code Useful, but still waiting for instructions	Own the workflow Deep judgment, but one person at a time	
Specialized reasoning	X No real business context or adaptive judgment	~ General coding help, not forecasting expertise	✓ But expensive and not always available	✓ Specialized reasoning for real model-building decisions
Fixes failures and retries	X Stops at the benchmark	X You still have to diagnose and steer	✓ If the person has time	✓ Auto-repair, retry, and continue in hours
Who can use it	ML teams Still needs technical setup	Technical users Prompting plus code review required	A data scientist Usually a bottleneck for everyone else	Data scientists and other professionals Sales, finance, growth,

WHY NOW · This wasn't possible 18 months ago. AI couldn't understand why it was wrong. Now it can – which means your team gets reliable, deployed models today, not after a hiring round or a six-month BI project.

SIMPLE, TRANSPARENT PRICING

From *\$20 a month* to enterprise rollout.

FOR INDIVIDUALS

\$20/month

A student, consultant, or analyst who needs answers from their data – without hiring anyone.

- Use it as much as you want
- Upload any spreadsheet
- Get a working model
- Download a clean report

MOST POPULAR

FOR TEAMS

\$500–2K/month

A small team that needs many predictions but can't afford to hire more data scientists.

- Whole team has access
- Connects to your systems
- Custom reports
- Email + chat support

FOR BIG COMPANIES

Custom

A large company that needs OctOpus inside their existing tools, with their security rules.

- Runs in your environment
- Dedicated support
- White-glove setup
- Pay per result

OUTCOME-BASED · For enterprise, we can price per result delivered – not per seat. You pay for outcomes, so the value is clear from day one. Start with a paid PoC and scale only once it proves itself on your data.

Built by a data scientist

Doudou BA

FOUNDER · ENTERPRISE AI & MACHINE LEARNING

OctOpus is built by someone who has shipped production AI inside a global enterprise – not a wrapper over someone else's model. Deep ML research, real enterprise deployment experience, and direct work with the kind of data and reliability requirements your team has. The reasoning that makes OctOpus trustworthy on your data comes from years of doing this work by hand.

EDUCATION

PhD Machine Learning · Charles University · MSc Software Engineering (Blekinge) · MSc Data Science (CZU)

RESEARCH

4 peer-reviewed publications across forecasting, hydrology, and ML systems

MAERSK

Built production AI and automation used in 130+ countries · the kind of reliability where minutes of downtime cost real money

DEPLOYMENT

Runs in your environment, against your warehouse, under your security rules – your data stays yours

AWARDS

Winner – Big Angels Day AI category · Winner – UNDP AI Hackathon · multiple Maersk Star Awards

ECOSYSTEM

NVIDIA Inception · Founders Inc · AI Sweden Ignite · AWS for Startups · Google Cloud Credits · Station F Fighters



Enterprise AI that's *actually reliable* — because it was built by someone who has been on the hook for it.

— THE EXPERTISE BEHIND OCTOPUS



PROOF OF CONCEPT · YOUR DATA · YOUR ENVIRONMENT

Start a proof of concept on *your own data* in weeks, not quarters.

HOW A POC WORKS

~2-4 WEEKS · LOW RISK

01 · CONNECT Pick one real problem We connect OctOpus to one dataset in your environment. Your data stays yours.	02 · BUILD See it solved OctOpus analyzes, builds the model and dashboards, and fills in a real report — on your data.	03 · DECIDE Keep it or walk away You see the result before you commit. Prove the value, then scale across your team.
---	---	---

START A POC →